

LLM Model Details

1. GPT-5.5

Model Provider: OpenAI

Model Type: Proprietary

Release Date: April 23, 2026

Knowledge Cutoff Date: December 1, 2025

Cost of Tokens:

- Input: \$5 / 1M tokens
- Output: \$30 / 1M tokens

Context Window: 1,050,000 tokens (1.05M)

Maximum Output Tokens: 128,000 tokens

Input Modalities: Text, Images, PDF

Output Modalities: Text

Use Cases:

1. Coding and debugging
2. Analyze data and spreadsheets
3. Research online
4. Creating reports, presentations, and documents

Practical Examples:

1. Build a complete web application from requirements, including code, tests, and documentation.
2. Analyze Excel and CSV files and generate charts, insights, and recommendations.
3. Draft emails, meeting summaries, and executive updates automatically.

2. Claude Opus 4.7

Model Provider: Anthropic

Model Type: Proprietary

Release Date: April 16, 2026

Knowledge Cutoff Date: January 2026

Cost of Tokens:

- Input: \$5 / 1M tokens
- Output: \$25 / 1M tokens

Context Window: 1,000,000 tokens (1M)

Maximum Output Tokens: 128,000 tokens

Input Modalities: Text, Images, PDF

Output Modalities: Text

Use Cases:

1. Complex coding and debugging
2. Agentic workflows
3. Research and knowledge work
4. Data and document analysis
5. Creating presentations and reports
6. Long-document summarization
7. Image and screenshot understanding
8. Scientific and technical reasoning

Practical Examples:

1. Analyze a large codebase and generate new features, refactoring suggestions, and tests.
2. Read hundreds of documents and produce a detailed literature review.

3. Gemini 3.1 Pro

Model Provider: Google DeepMind

Model Type: Proprietary

Release Date: February 19, 2026

Knowledge Cutoff Date: January 2025

Cost of Tokens:

- Input: \$2 / 1M tokens
- Output: \$12 / 1M tokens

Context Window: 1,000,000 tokens (1M)

Maximum Output Tokens: 65,536 tokens

Input Modalities: Text, Images, Audio, Video, PDF

Output Modalities: Text

Use Cases:

1. Complex reasoning and problem solving
2. Writing and debugging code
3. Long-document analysis and summarization
4. Multimodal analysis (images, videos, and PDFs)
5. Research and knowledge synthesis

6. Data analysis and report generation
7. Agentic workflows and tool use
8. Scientific and technical tasks

Practical Examples:

1. Build a complete web application from a single prompt, including frontend, backend, tests, and deployment scripts.
2. Analyze PDFs, images, audio, and videos together to produce a detailed report.
3. Gather information from multiple sources and create a citation-backed research report.

4. DeepSeek-V4 (Flash and Pro)

Model Provider: DeepSeek

Model Type: Open-source

Release Date: April 24, 2026

Knowledge Cutoff Date: Not officially disclosed

Cost of Tokens:

- Input: \$1.74 / 1M tokens, Output: \$3.48 / 1M tokens (**Pro**)
- Input: \$0.14 / 1M tokens, Output: \$0.28 / 1M tokens (**Flash**)

Context Window: 1,000,000 tokens (1M)

Maximum Output Tokens: 384,000 Tokens

Input Modalities: : Text, Images, Documents

Output Modalities: Text

Use Cases:

1. Coding and Debugging
2. Long-document analysis and summarization
3. Research and knowledge synthesis
4. Agentic workflows and tool use
5. Data analysis and report generation
6. Scientific and technical reasoning
7. Large-context retrieval and planning
8. Cost-efficient self-hosting and fine-tuning

Practical Examples:

1. Private AI Deployment- Self-host a secure assistant for internal company use.
2. Multi-Agent Systems - Coordinate specialized agents for coding, research, and automation.

5. Grok-4.20

Model Provider: xAI

Model Type: Proprietary

Release Date: March 9, 2026

Knowledge Cutoff Date: September 1, 2025

Cost of Tokens:

- Input: \$1.25 / 1M tokens
- Output: \$2.50 / 1M tokens

Context Window: 2,000,000 tokens

Maximum Output Tokens: 256,000 tokens

Input Modalities: Text, Images

Output Modalities: Text

Use Cases:

1. Advanced reasoning
2. Multi-agent research
3. Large-codebase analysis
4. Scientific research
5. Real-time information synthesis

Practical Examples:

1. Build a market intelligence assistant that monitors news, social posts, and industry updates in real time.
2. Track competitor announcements and summarize strategic insights automatically

6. Qwen3

Model Provider: Qwen

Model Type: Open-source

Release Date: April 28, 2025

Knowledge Cutoff: March 31, 2025

Cost Of Tokens: Varies by provider

Context Window: 128,000 tokens

Maximum Output Tokens: Varies by deployment and model size

Input Modalities: Text

Output Modalities: Text

Use Cases:

1. Writing and debugging code
2. Research and summarization
3. Multilingual translation
4. Data analysis
5. Document generation
6. Agentic workflows and tool use
7. Mathematical reasoning
8. Local/private AI deployments
9. Supports 119 languages and dialects

Practical Examples:

1. Build a multilingual customer support assistant.
2. Run a secure internal AI assistant on local infrastructure.
3. Summarize a 100-page PDF into key insights.

7. Mistral Large 2

Model Provider: Mistral AI

Model Type: Proprietary

Model Release Date: July 24, 2024

Knowledge Cutoff Date: Not officially disclosed

Cost Of Tokens:

- Input: ~\$2 / 1M tokens
- Output: ~\$6 / 1M tokens

Context Window: 128,000 tokens

Maximum Output Tokens: 8,000–16,000 tokens

Input Modalities: Text

Output Modalities: Text

Use Cases:

1. Complex reasoning
2. Writing and debugging code
3. Multilingual translation
4. Long-document summarization
5. Structured JSON generation

Practical Examples:

1. Write a two-page project status report.
2. Debug a Python script and explain the root cause.
3. Extract invoice details into JSON.

8. Llama 4 (Scout, Maverick)

Model Provider: Meta AI

Model Type: Open source

Release Date: April 5, 2025

Knowledge Cutoff Date: August 2024

Cost of Tokens :

- Scout: Input \$0.08 / 1M; Output \$0.30 / 1M
- Maverick: Input \$0.15–\$0.20 / 1M; Output \$0.60 / 1M

Context Window: Scout: 10M tokens; Maverick: 1M tokens

Maximum Output Tokens: 16K tokens

Input Modalities: Text, Images

Output Modalities: Text

Use Cases:

1. Coding and debugging
2. Multimodal document analysis
3. Knowledge assistants over private data
4. Self-hosted enterprise AI

Practical Examples:

1. Build a secure internal chatbot.
2. Generate and review code for software projects.
3. Translate content across multiple languages.

9. Nemotron-3

Model Provider: NVIDIA

Model Type: Open-weight

Release Date: December 15, 2025

Knowledge Cutoff Date: November 28, 2025

Cost Of Tokens:

- Nano: Input ~\$0.05–\$0.20 / 1M; Output ~\$0.10–\$0.40 / 1M
- Super: Input ~\$0.10–\$0.30 / 1M; Output ~\$0.50–\$0.80 / 1M

Context Window: 1,000,000 tokens (1M)

Maximum Output Tokens: 16K–32K tokens

Input Modalities: Text

Output Modalities: Text

Use Cases:

1. Agentic AI and multi-agent systems
2. Coding and debugging
3. Long-document analysis
4. Retrieval-Augmented Generation (RAG)
5. Enterprise knowledge assistants
6. Workflow automation
7. Research and summarization

Practical Examples:

1. RAG Systems - Create a knowledge assistant that retrieves and cites information from company data.
2. Long-Document Analysis - Analyze contracts, research papers, and technical manuals with up to 1M tokens of context.
3. Scientific Computing - Assist with simulations, experiment design, and technical reasoning.

10. Kimi K2.6

Model Provider: Moonshot AI

ModelType: Open source (open-weight)

Release Date: April 20, 2026

Knowledge Cutoff Date: Not officially disclosed

Cost of Tokens:

- Input: ~\$0.60–\$0.75 / 1M tokens;
- Output: ~\$3.50–\$4.00 / 1M tokens

Context Window: 256K tokens

Maximum Output Tokens: 65,536 tokens

Input Modalities: Text, Images, Video

Output Modalities: Text

Use Cases:

1. Long-horizon coding and debugging
2. Multi-agent workflows
3. Document and spreadsheet analysis
4. Research and report generation
5. Vision and video understanding
6. Tool calling and workflow automation
7. Self-hosted enterprise AI assistants

Practical Examples:

1. Self-Healing Automation - Suggest locator fixes when elements change.
2. Test Script Generation - Generate Playwright or Selenium scripts from manual test cases.
3. Test Maintenance - Update hundreds of scripts after UI or API changes.
4. Regression Suite Optimization - Detect redundant tests and recommend a smaller high-value suite.
5. Performance Test Creation - Generate load-testing scripts and assertions.